

P. H. D

Churn Classification

By

Karthik Reddy Mathuru

1404

Batch-31

Problem Statement:

To predict the telecom customers who are likely to exit the contract and also to generate patterns of Churn and non-churn to assist the management to take appropriate decisions to limit churn.

Most telecom companies suffer from voluntary churn. Churn rate has strong impact on the life time value of the customer because it affects the length of service and the future revenue of the company. It is estimated that 75 percent of the 17 to 20 million subscribers signing up with a new wireless carrier every year are coming from another wireless provider, which means they are churners.

State of the art:

Telecom companies are battling to attract each other's customers while retaining their own. Thus, Customer churn reduction is the central concern of most telecom companies as switching costs to the customer are low and acquisition cost to the company is high. Churn reduces profitability as it means potential loss of future revenue and also losing the invested costs of acquisition.

Ways to Reduce Customer Churn:

- Consistently Exceed Customers' Expectations

The most fundamental way to decrease your churn rate is by keeping your customers happy. While you definitely want to avoid letting them down, you have to look for areas to go over and above your customer's expectations and delight them. Failing to deliver on a promise is one of the fastest ways to lose a customer, and many companies say that dissatisfaction and unmet expectations are among the top reasons for client churn.

- Provide Awesome Customer Service

This one should go without saying, but if you've ever spent half an hour listening to hold muzak waiting for a disinterested, incompetent customer service rep to "assist you," you'll know that some companies simply don't put enough effort into customer service.

- **Create Switching Costs:**

Switching costs are any cost that a customer incurs by trading one product or service for another. Higher switching costs naturally reduce churn by reducing the likelihood that a customer will switch to a substitute product instead of returning to your brand.

Methods:

Logistic Regression:

Logistic regression is a statistical method for analyzing a dataset in which there are one or more independent variables that determine an outcome. The outcome is measured with a dichotomous variable (in which there are only two possible outcomes).

- We use ROC (Receiver Operating Characteristics) curve to set a threshold value which minimizes the false positive rate and maximizes the true positive rate.

Decision Trees:

Decision Tree is a type of supervised learning algorithm (having a pre-defined target variable) that is mostly used in classification problems. It works for both categorical and continuous input and output variables. In this technique, we split the population or sample into two or more homogeneous sets using greedy approach.

- Gini Index – (Classification and Regression Trees)
- Entropy – (C 5.0)

The above gives the impurity of the node. By these values we choose the node.

Random Forests:

Random forests are an ensemble learning method for classification, regression and other tasks, that operate by constructing a multitude of decision trees at training time and outputting the class that is the mode of the classes (classification) or mean prediction (regression) of the individual trees.

- It performs bagging technique (divides dataset into samples with replacement and build a decision tree for each sample)
- Random forest takes random number of features while performing bagging.
- It reduces the variance, there is less chance of overfitting.

Xg- Boost:

XGBoost is an open-source software library which provides the gradient boosting framework for C++, Java, Python, R and Julia. It works on Linux, Windows and macOS. From the project description, it aims to provide a "Scalable, Portable and Distributed Gradient Boosting (GBM, GBRT, GBDT) Library".

- XGBoost is extreme gradient boosting technique
- It is a boosting technique (builds a decision tree and add more weight to wrongly classified sample, again built a tree on it and so on till n trees)
- It reduces both bias and variance.

SVM:

“Support Vector Machine” (SVM) is a supervised machine learning algorithm which can be used for both classification or regression challenges. However, it is mostly used in classification problems. It finds a max margin hyper plane two separate classes.

- In SVM, we have many kernel – radial, polynomial, linear etc.
- It performs quadratic optimization and kernel trick
- It also handles dual problem.
- It can handle class imbalance.

Stacking:

Stacking (also called meta ensembling) is a model ensembling technique used to combine information from multiple predictive models to generate a new model. Often times the stacked model (also called 2nd-level model) will outperform each of the individual models due its smoothing nature and ability to highlight each base model where it performs best and discredit each base model where it performs poorly.

MLP:

A multilayer perceptron (MLP) is a class of feedforward artificial neural network. An MLP consists of at least three layers of nodes. Except for the input nodes, each node is a neuron that uses a nonlinear activation function. MLP utilizes a supervised learning technique called backpropagation for training. Its multiple layers and non-linear activation distinguish MLP from a linear perceptron. It can distinguish data that is not linearly separable. Multilayer perceptron's are sometimes colloquially referred to as "vanilla" neural networks, especially when they have a single hidden layer

Data:

The dataset depicts the details of the telecom customer, like, services that each customer has signed for, customer account information, and their demographic info etc.

- Attributes – 25 (Numerical – 3, Categorical – 20, Date – 2)
- Records – 5925

Attribute Information:

Demographics Data:

- HouseholdID: Each Household id
- Country: Country (For this attribute, missing values are denoted as “?”)
- State: State (For this attribute, missing values are denoted as “?”)
- Retired: Whether retired
- HasPartner: Demographic information - whether the customer has partner (1-Yes; 2-No)
- HasDependents: Demographic information - whether the customer has dependents (1-Yes; 2-No)
- Education: Education qualification
- Gender: Demographic information – gender

Account Information:

- CustomerID: CustomerID
- Base Charges: Customer account information (Charges for Base plan)
- DOC: Date of data collection
- Total Charges: Customer account information (Total). (For this attribute, missing values are denoted as “MISSINGVAL” also)

- DOE: Date of entry as customer
- Electronic Billing: Customer account information - whether electronic billing
- Contract Type: Contract type (For this attribute, missing values are denoted as "NA")
- Payment Method: payment method

Data of ServicedOptedFor:

- CustomerID: CustomerID
- TypeOfService: Service signed for
- ServiceDetails:

Churn Data:

- CustomerID: Customer ID
- Churn: Whether the customer churns (Target)

Pre-processing:

- The datasets have the details of the customer - services that each customer has signed for, customer account information, and their demographic info.
- Data is aggregated from four datasets. Customer ID is the primary key
- Data in services have long format, it is made to wide format using pivot table(python).
- There are 42 missing values in total. We have two columns which have no variance (Country, State) which has 13 missing values, these columns are removed.
- Missing values are very less. So, omitted the records with NA.

Feature Engineering:

- DOC (Date of Collection)
- DOE (Date of Entry)

From above two columns we get the age (Life time) of the customer by difference between them.

$$\text{AGE} = \text{DOC} - \text{DOE}$$

Standardization:

Age, Base Charges, Total Charges are the numeric attributes in data with different scale. We use standardization techniques to make them in same scale.

Correlation:


By the plot, we know that there is a high positive correlation between Age (attribute is derived from DOC and DOE) and Total charges.



Analysis:


➤ Uni – Variate:

Contract Type	Count of Contract Type
One year	1,150
Two years	1,364
Month-to-month	2,779




Count of Contract Type for each Contract Type. The marks are labeled by count of Contract Type. The view is filtered on Contract Type, which excludes NA.

Device Protection	Count of Device Protection
Yes	1,889
No	3,409




Count of Device Protection for each Device Protection. The marks are labeled by count of Device Protection.

Education	Count of Education
Other	41
Professional Qualification	1,186
Masters	1,265
Graduation	1,370
Highschool or below	1,426




Count of Education for each Education. The marks are labeled by count of Education. The view is filtered on Education, which excludes Null.

Electronic Billing	Count of Electronic Billing
No	2,178
Yes	3,120




Count of Electronic Billing for each Electronic Billing. The marks are labeled by count of Electronic Billing.

Gender	Count of Gender
Female	2,638
Male	2,656




Count of Gender for each Gender. The marks are labeled by count of Gender. The view is filtered on Gender, which excludes Null.

Internet Service Category	Count of Internet Service Category
No	1,156
DSL	1,829
Fiber optic	2,313




Count of Internet Service Category for each Internet Service Category. The marks are labeled by count of Internet Service Category.

Multiple Lines	Count of Multiple Lines
Yes	2,301
No	2,997




Count of Multiple Lines for each Multiple Lines. The marks are labeled by count of Multiple Lines.

Online Backup	Count of Online Backup
Yes	1,872
No	3,426




Count of Online Backup for each Online Backup. The marks are labeled by count of Online Backup.

Online Security	Count of Online Security
Yes	1,573
No	3,725



Count of Online Security for each Online Security. The marks are labeled by count of Online Security.

Payment Method	Count of Payment Method
Mailed check	1,033
Bank transfer (automatic)	1,149
Credit card (automatic)	1,530
Electronic check	1,586

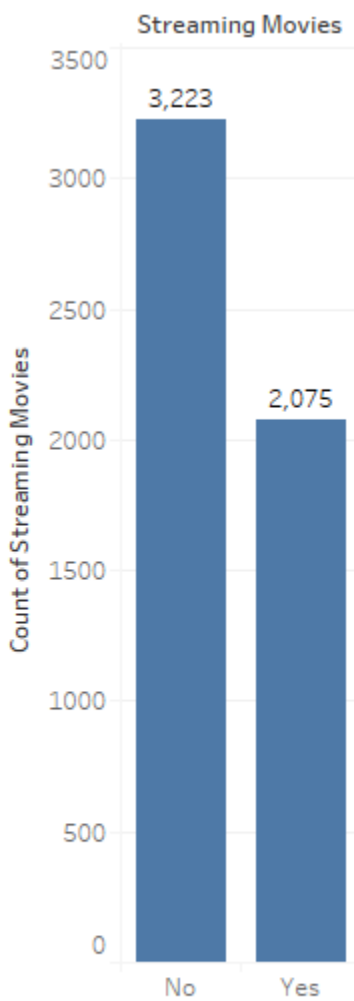


Count of Payment Method for each Payment Method. The marks are labeled by count of Payment Method.

Streaming Movies Count of Streaming
Movies

Yes 2,075

No 3,223




Count of Streaming Movies for
each Streaming Movies. The
marks are labeled by count of
Streaming Movies.

Streaming Television Count of Streaming
Television


Yes 2,075

No 3,223




Count of Streaming Television
for each Streaming Television.
The marks are labeled by count
of Streaming Television.

Technical Support	Count of Technical Support
Yes	1,605
No	3,693



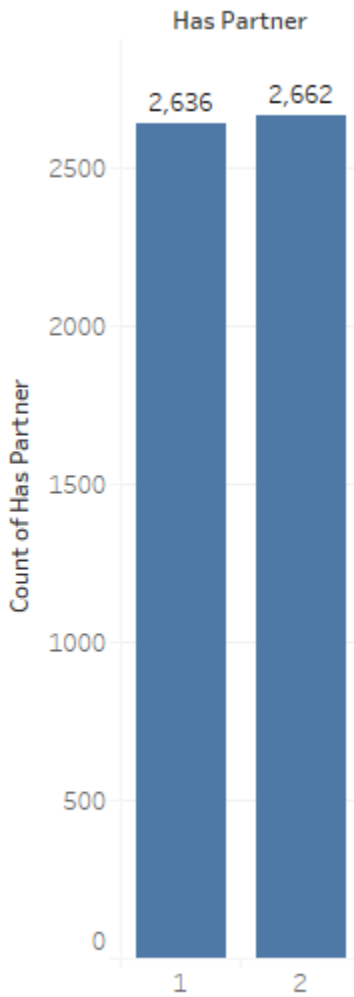
Count of Technical Support for each Technical Support. The marks are labeled by count of Technical Support.

Has Dependents	Count of Has Dependents
1	1,638
2	3,660




Count of Has Dependents for each Has Dependents. The marks are labeled by count of Has Dependents.

Has Partner	Count of Has Partner
1	2,636
2	2,662




Count of Has Partner for each Has Partner. The marks are labeled by count of Has Partner.

Has Phone Service	Count of Has Phone Service
0	506
1	4,792




Count of Has Phone Service for each Has Phone Service.

Age:




The trend of count of Age for Age (bin).

Base Charges:



The trend of count of Base Charges for Base Charges (bin).

Total Charges:



The trend of count of Total Charges for Total Charges (bin).

Result:

Error metric is Recall

- Recall = True Positive/ Total Actual Positive
- We should predict telecom customers who are likely to churn. Recall of class 'Yes' gives us customers who are likely to churn.

➤ Decision Tree:

- Gini index

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.83	0.83	790
1	0.48	0.48	0.48	264
avg / total	0.74	0.74	0.74	1054

Accuracy 0.740037950664

- Entropy:

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.83	0.83	790
1	0.49	0.50	0.49	264
avg / total	0.75	0.74	0.74	1054

Accuracy : 0.743833017078

➤ Random Forest:

Classifaction Report :

	precision	recall	f1-score	support
0	0.91	0.76	0.83	790
1	0.52	0.77	0.62	264
avg / total	0.81	0.76	0.78	1054

Accuracy : 0.76375711575

➤ Xg Boost

Classification Report :

	precision	recall	f1-score	support
0	0.85	0.91	0.88	790
1	0.66	0.51	0.58	264
avg / total	0.80	0.81	0.80	1054

Accuracy : 0.812144212524

➤ SVM

Classification Report :

	precision	recall	f1-score	support
0	0.89	0.82	0.85	790
1	0.56	0.70	0.63	264
avg / total	0.81	0.79	0.80	1054


Accuracy : 0.789373814042

➤ Logistic regression

Classification Report :

	precision	recall	f1-score	support
0	0.90	0.75	0.82	790
1	0.51	0.76	0.61	264
avg / total	0.80	0.75	0.77	1054

Accuracy : 0.753320683112



➤ Stacking:

Classification Report :

	precision	recall	f1-score	support
0	0.88	0.84	0.86	790
1	0.59	0.67	0.63	264
avg / total	0.81	0.80	0.80	1054

Accuracy : 0.799810246679

➤ MLP:

Classification Report :

	precision	recall	f1-score	support
0	0.84	0.87	0.86	790
1	0.57	0.52	0.54	264
avg / total	0.78	0.78	0.78	1054

Accuracy : 0.781783681214

Patterns:

Decision trees are used to extract the hidden patterns, decision trees give rules which are supported by confidence, lift. C5.0 is used to extract the patterns.

Appendices:

```
# In[1]:
```

```
# import modules
```

```
import pandas as pd
```

```
import numpy as np
```

```
# ##### Loading the data
```

```
# In[2]:
```

```
# Train
```

```
train = pd.read_csv('TrainData\Train.csv')
```

```
accinfo = pd.read_csv('TrainData\Train_Accountinfo.csv',na_values=['NA','MISSINGVAL'])
```

```
demographics = pd.read_csv('TrainData\Train_Demographics.csv',na_values='?')
```

```
services = pd.read_csv('TrainData\Train_ServicesOptedFor.csv')
```

```
# In[3]:
```

```
# Test
```

```
test = pd.read_csv('TestData\Test.csv')
```

```
taccinfo = pd.read_csv('TestData\Test_Accountinfo.csv',na_values=['NA','MISSINGVAL'])
```

```
tdemographics = pd.read_csv('TestData\Test_Demographics.csv',na_values='')
tservices = pd.read_csv('TestData\Test_ServicesOptedFor.csv')
```

```
# #### Dimensions
```

```
# In[4]:
```

```
# Shapes of the train data
print(train.shape)
print(accinfo.shape)
print(demographics.shape)
print(services.shape)
```

```
# In[5]:
```

```
# Shapes of test data
print(test.shape)
print(taccinfo.shape)
print(tdemographics.shape)
print(tservices.shape)
```

```
# From the data, we know that we have CustomerID as PrimaryKey
```

```
# In[6]:
```

```
# unique columns in datasets
print(len(train.CustomerID.unique()))
print(len(accinfo.CustomerID.unique()))
print(len(demographics.HouseholdID.unique()))
print(len(services.CustomerID.unique()))
```

```
# ##### Merging the data frames by CustomerID
```

```
# In[7]:
```

```
# Train
```

```
n1 = pd.merge(train,accinfo,on='CustomerID')
```

```
# In[8]:
```

```
# We have HouseholdID in demographics which is CustomerID so renaming the column.
```

```
demographics.columns = ['CustomerID', 'Country', 'State', 'Retired', 'HasPartner',
    'HasDependents', 'Education', 'Gender']
```

```
# In[9]:
```

```
new = pd.merge(n1,demographics,on='CustomerID')
```

```
# In[10]:
```

```
# Dimensions
```

```
new.shape
```

```
# In[11]:
```

```
# Test
```

```
tdemographics.columns = ['CustomerID', 'Country', 'State', 'Retired', 'HasPartner',  
    'HasDependents', 'Education', 'Gender']
```

```
# In[12]:
```

```
tnew = pd.merge(taccinfo,tdemographics,on='CustomerID')
```

```
# In[13]:
```

```
tnew.shape
```

```
# ** Services ** have shape (15921, 3)
```

```
# In[14]:
```

```
# count - categories
```

```
services.SeviceDetails.value_counts()
```



```
# In[15]:
```

```
# Changing the attribute types to category
```

```
services.ServiceDetails = services.ServiceDetails.astype('category')
```

```
# In[16]:
```

```
# Test
```

```
tservices.ServiceDetails = tservices.ServiceDetails.astype('category')
```

```
# * Dataframe is in long format, we should make it into wide format
```

```
# In[17]:
```

```
# Label Encoding
```

```
services.ServiceDetails = services.ServiceDetails.cat.codes
```

```
services.ServiceDetails.value_counts()
```

```
# In[18]:
```

```
tservices.ServiceDetails = tservices.ServiceDetails.cat.codes
```

```
tservices.ServiceDetails.value_counts()
```

```
# In[19]:
```

```
# creating a multilevel index with values as serviceDetails
```

```
n = services.pivot_table(index="CustomerID", columns="TypeOfService", values = "SeviceDetails")
```

```
# In[20]:
```

```
n.head()
```

```
# In[21]:
```

```
ser = np.matrix(n)
```

```
# In[22]:
```

```
ser = pd.DataFrame(ser, columns=n.columns)
```

```
ser.head()
```

```
# In[23]:
```

```
tn = tservices.pivot_table(index="CustomerID", columns="TypeOfService", values = "SeviceDetails")
```

```
# In[24]:
```

```
tser = np.matrix(tn)
```

```
# In[25]:
```

```
tser = pd.DataFrame(tser,columns=tn.columns)
```

```
tser.head()
```

```
# In[26]:
```

```
# concatenating the data frames
```

```
final = pd.concat([new,ser],1) # train
```

```
tfinal = pd.concat([tnew,tser],1) # test
```

```
# In[27]:
```

```
final.head() # view
```

```
# In[28]:
```

```
tfinal.head()
```

```
# #### Missing Values
```

```
# In[29]:
```

```
print(final.isnull().sum())  
print('\n Total:\t\t', final.isnull().any(1).sum())
```

```
# In[30]:
```

```
# Drop the columns with no variance and unique values
```

```
final = final.drop(['CustomerID','Country','State'],1)  
tfinal = tfinal.drop(['CustomerID','Country','State'],1)
```

```
# In[31]:
```

```
print(final.isnull().sum())  
print('\n Total:\t\t', final.isnull().any(1).sum())
```

```
# In[32]:
```

```
# Missing values are missing at random, so there is no pattern in it.
```

```
final[final.isnull().any(1)]
```

```
# In[33]:
```

```
final = final.dropna() # Drop missing values
```

```
# In[34]:
```

```
final.to_csv('f.csv',index=False)
```

```
final = pd.read_csv('f.csv')
```

```
# In[35]:
```

```
final.head()
```

```
# ##### Type Casting
```

```
# In[36]:
```

```
num = [ 'BaseCharges', 'TotalCharges' ]
```

```
date = ['DOC', 'DOE']
```

```
cat = ['Churn','ElectronicBilling','ContractType', 'PaymentMethod',  
      'Retired', 'HasPartner', 'HasDependents', 'Education',  
      'Gender','DeviceProtection', 'HasPhoneService',  
      'InternetServiceCategory', 'MultipleLines', 'OnlineBackup',  
      'OnlineSecurity', 'StreamingMovies', 'StreamingTelevision',  
      'TechnicalSupport']
```

```
# ##### Categorical
```

```
# In[37]:
```

```
# Changing attributes to category
```

```
for i in cat:
```

```
    final[i] = final[i].astype('category')
```

```
# In[38]:
```

```
# Label Encoding for categorical attributes
```

```
for i in cat:
```

```
    final[i] = final[i].cat.codes
```

```
for i in cat:
```

```
    final[i] = final[i].astype('category')
```

```
# In[39]:
```

```
# Test
```

```
for i in cat[1:]:
```

```
    tfinal[i] = tfinal[i].astype('category')
```

```
# In[40]:
```

```
for i in cat[1:]:  
    tfinal[i] = tfinal[i].cat.codes  
for i in cat[1:]:  
    tfinal[i] = tfinal[i].astype('category')
```

```
# In[41]:
```

```
final.to_csv('fin.csv')# write data
```

```
# In[42]:
```

```
final.head()
```

```
# ### UniVariate Analysis:
```

```
# In[43]:
```

```
import matplotlib.pyplot as plt  
get_ipython().magic('matplotlib inline')
```

```
# In[44]:
```

```
# Plots for categorical attributes
```

```
for i in cat:
```

```
pd.crosstab(final[i],columns='count').plot(kind = 'bar')  
plt.savefig(i)
```

```
# ### Numerical
```

```
# In[45]:
```

```
for i in num:  
    final[i] = final[i].astype('int')  
# Base Charges plot  
plt.figure()  
final.BaseCharges.hist(grid = False)  
plt.xlabel('BaseCharges')  
plt.ylabel('frequency')  
plt.savefig('base_hist.png')
```

```
# In[46]:
```

```
# Total Charges plot  
plt.figure()  
final.TotalCharges.hist()  
plt.xlabel('TotalCharges')  
plt.ylabel('frequency')  
plt.savefig('tot_hist.png')
```



```
# In[47]:
```

```
# Box plots
```

```
for i in num:
```

```
    final.boxplot(i)
```

```
    plt.show()
```

```
    plt.savefig(i)
```

```
# ### Date
```

```
# In[48]:
```

```
import datetime
```

```
# In[49]:
```

```
# Date time format
```

```
final.DOC = pd.to_datetime(final.DOC)
```

```
tfinal.DOC = pd.to_datetime(tfinal.DOC)
```

```
# In[50]:
```

```
final.DOE = pd.to_datetime(final.DOE)
```

```
tfinal.DOE = pd.to_datetime(tfinal.DOE)
```

```
# #### Feature Engineering
```

```
# In[51]:
```

```
# Age is the lifetime of the customer
```

```
final['Age'] = final.DOC - final.DOE
```

```
tfinal['Age'] = tfinal.DOC - tfinal.DOE
```

```
# In[52]:
```

```
# Converting to int, models does not accept timedelta
```

```
final.Age = (final.Age / np.timedelta64(1, 'D')).astype(int)
```

```
tfinal.Age = (tfinal.Age / np.timedelta64(1, 'D')).astype(int)
```

```
# In[53]:
```

```
tfinal.head()
```

```
# ** correlation **
```

```
# In[54]:
```

```
import seaborn as sns
```

```
sns.heatmap(final[['BaseCharges','Age','TotalCharges']].corr())
```

```
# #### Standardization
```

```
# In[55]:
```

```
# Numerical variables
```

```
num_var = final[['BaseCharges','Age' ]]
```

```
num_var.shape
```

```
tnum_var = tfinal[['BaseCharges', 'Age' ]]
```

```
tnum_var.shape
```

```
# In[56]:
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler()
```

```
# In[57]:
```

```
num_scaled = pd.DataFrame(scaler.fit_transform(num_var),columns=num_var.columns)
```

```
num_scaled.head()
```

```
# In[58]:
```

```
tnum_var.isnull().sum()
```

```
# In[59]:
```

```
tnum_scaled = pd.DataFrame(scaler.fit_transform(tnum_var),columns=tnum_var.columns)
```

```
# In[60]:
```

```
tnum_scaled.shape
```

```
# In[61]:
```

```
# Dropping all the columns as we already have them
```

```
final = final.drop(['TotalCharges','BaseCharges','Age'],1)
```

```
# In[62]:
```

```
tfinal = tfinal.drop(['TotalCharges','BaseCharges','Age'],1)
```

```
# In[63]:
```

```
final = pd.concat([tnum_scaled,final],1) # concatenation
```

```
# In[64]:
```

```
final.head()
```

```
# In[65]:
```

```
tfinal = pd.concat([tnum_scaled,tfinal],1)
```

```
# In[66]:
```

```
tfinal.head()
```

```
# In[67]:
```

```
final.columns # Columns
```

```
# In[68]:
```

```
final.shape #dimensions
```

```
# In[69]:
```

```
final.to_csv('tab.csv')
```

```
# In[70]:
```

```
te_x = tfinal.drop(['DOC','DOE'],1) # derived new column from DOC and DOE, so we drop them
```

```
# In[71]:
```

```
te_x.head()
```

```
# In[72]:
```

```
y = final.Churn # Target
```

```
x = final.drop(['Churn','DOC','DOE'],1)
```

```
# In[73]:
```

```
x.head()
```

```
# In[74]:
```

```
# Percentage of target variable
```

```
print('0 :',(y.value_counts()[0]/y.value_counts().sum())*100)
```

```
print('1 :',(y.value_counts()[1]/y.value_counts().sum())*100)
```

```
# ### Train Test Split
```

```
# In[75]:
```

```
from sklearn.model_selection import train_test_split
```

```
# In[76]:
```

```
xtrain,xtest,ytrain,ytest = train_test_split(x,y,test_size = 0.2,stratify = y)
```

```
# ### Model Building:
```

```
# **Decision tree :**
```

```
#
```

Decision Tree is a type of supervised learning algorithm (having a pre-defined target variable) that is mostly used in classification problems. It works for both categorical and continuous input and output variables. In this technique, we split the population or sample into two or more homogeneous sets using greedy approach.

```
# In[77]:
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
from sklearn import metrics
```

```
# In[78]:
```

```
# gini index - CART(Classification and Refression Tree)
```

```
tree = DecisionTreeClassifier(class_weight='balanced')
```

```
# In[79]:
```

```
tree.fit(xtrain,ytrain) # Train the model
```

```
# In[80]:
```

```
tree.score(xtrain,ytrain) #Accuracy
```

```
# In[81]:
```

```
print('\n Classification Report:\n\n',metrics.classification_report(ytest,tree.predict(xtest)))
```

```
print('Accuracy',metrics.accuracy_score(ytest,tree.predict(xtest)))
```

```
# In[82]:
```

```
# Entropy - C5.0
```

```
entropy = DecisionTreeClassifier(criterion='entropy',class_weight='balanced') # entropy
```



```
# In[83]:
```

```
entropy.fit(xtrain,ytrain) #Train the model
```

```
# In[84]:
```

```
entropy.score(xtest,ytest) # accuracy
```

```
# In[85]:
```

```
print('\n Classification Report:\n\n',metrics.classification_report(ytest,entropy.predict(xtest)))
```

```
print('Accuracy :',metrics.accuracy_score(ytest,entropy.predict(xtest)))
```

```
# **Random Forest:**
```

```
#
```

```
# Random Forests are an ensemble learning method for classification, regression and other tasks, that  
operate by constructing a multiple decision trees at training time and outputting the class that is the  
mode of the classes for classification and mean prediction for regression.
```

```
# In[86]:
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
# In[87]:
```

```
r = RandomForestClassifier(n_estimators=100,n_jobs=-1,oob_score=True,  
                           class_weight='balanced',max_depth=14,random_state = 20,max_leaf_nodes =  
29,min_samples_leaf = 6)
```

```
# In[88]:
```

```
rf_exp = r.fit(xtrain,ytrain)
```

```
# In[89]:
```

```
r.score(xtrain,ytrain)
```

```
# In[90]:
```

```
print('\n Classification Report : \n\n',metrics.classification_report(ytest,r.predict(xtest)))
```

```
print('Accuracy :',metrics.accuracy_score(ytest,r.predict(xtest)))
```

```
# In[91]:
```

```
# Tuning
```

```
from sklearn.grid_search import RandomizedSearchCV
```

```
# In[92]:
```

```
param_grid = dict(max_depth=list(range(1, 20)),min_samples_leaf = list(range(1,12)),  
                  max_leaf_nodes=list(range(2,30)))
```

```
# In[93]:
```

```
grid = RandomizedSearchCV(r, param_grid, cv=10, scoring='accuracy',n_jobs= -1,n_iter=20)
```

```
# In[94]:
```

```
grid.fit(xtrain,ytrain)
```

```
# In[95]:
```

```
grid.best_params_
```

```
# ** Random Forest : Feature Importance **
```

```
# In[96]:
```

```
sorted(list(zip(r.feature_importances_,xtrain.columns)))
```

```
# In[97]:
```

```
### Drop the columns which are not important
```

```
rxtrain = xtrain.drop(['Retired', 'HasPartner', 'Gender',  
                      'HasPhoneService', 'DeviceProtection', 'MultipleLines',  
                      'HasDependents', 'StreamingTelevision'], 1)  
rxtest = xtest.drop(['Retired', 'HasPartner', 'Gender',  
                    'HasPhoneService', 'DeviceProtection', 'MultipleLines',  
                    'HasDependents', 'StreamingTelevision'], 1)
```

```
# In[98]:
```

```
rxtrain.columns
```

```
# In[99]:
```

```
rte_x = te_x.drop(['Retired', 'HasPartner', 'Gender',  
                  'HasPhoneService', 'DeviceProtection', 'MultipleLines',  
                  'HasDependents', 'StreamingTelevision'], 1) # Test
```

```
# In[100]:
```

```
rte_x.columns
```

```
# ** Xg Boost: **
```

```
# In[101]:
```

```
from xgboost import XGBClassifier
```

```
# In[102]:
```

```
xg = XGBClassifier(n_jobs=-1,n_estimators= 500,booster='gblinear',max_delta_step= 2)
```

```
# In[103]:
```

```
xg.fit(np.matrix(xtrain),ytrain)
```

```
# In[104]:
```

```
xg.score(np.matrix(xtrain),ytrain) # Predictions
```

```
# In[105]:
```

```
print('\n Classifaction Report : \n\n',metrics.classification_report(ytest,xg.predict(np.matrix(xtest)))) #  
Predictions
```

```
print('Accuracy :',metrics.accuracy_score(ytest,xg.predict(np.matrix(xtest))))
```

```
# ** SVM (Support Vector Machine ):**
```

```
# In[106]:
```

```
from sklearn.svm import LinearSVC
```

```
s = LinearSVC(class_weight='balanced',C = 10)
```

```
# In[107]:
```

```
s.fit(xtrain,ytrain)
```

```
# In[108]:
```

```
print('\n Classification Report : \n\n',metrics.classification_report(ytest,s.predict(xtest))) # Predictions
```

```
print('Accuracy :',metrics.accuracy_score(ytest,s.predict(xtest)))
```

```
# **Logistic Regression : **
```

```
# In[109]:
```

```
from sklearn.linear_model import LogisticRegressionCV
```

```
logreg = LogisticRegressionCV(class_weight='balanced',random_state=20,cv =  
10,scoring='accuracy',Cs=100)
```

```
# In[110]:
```

```
logreg.fit(xtrain,ytrain)
```

```
# In[111]:
```

```
import scikitplot as skplt
```

```
skplt.metrics.plot_roc_curve(ytrain,logreg.predict_proba(xtrain))
```

```
# In[112]:
```

```
print('\n Classification Report : \n\n',metrics.classification_report(ytest,logreg.predict(xtest)))
```

```
print('Accuracy :',metrics.accuracy_score(ytest,logreg.predict(xtest)))
```

```
# ##### Stacking
```

```
# In[113]:
```

```
# Train Predictions
```

```
sx = s.predict(np.matrix(xtrain))
```

```
sl = logreg.predict(xtrain)
```

```
sr = r.predict(xtrain)
```

```
# In[114]:
```

```
# val Predictions
```

```
stx = s.predict(np.matrix(xtest))
```

```
strr = r.predict(xtest)
```

```
stl = logreg.predict(xtest)
```

```
# In[115]:
```

```
# test Predictions
```

```
stxt = s.predict(np.matrix(te_x))
```

```
strrt = r.predict(te_x)
```

```
stlt = logreg.predict(te_x)
```

```
# In[116]:
```

```
# Train
```

```
stack_train = pd.concat([pd.DataFrame(sx),pd.DataFrame(sr),pd.DataFrame(sl)],1)
```

```
# In[117]:
```

```
# val
```

```
stack_test = pd.concat([pd.DataFrame(stx),pd.DataFrame(strr),pd.DataFrame(stl)],1)
```



```
# In[118]:
```

```
# test
```

```
stack_t = pd.concat([pd.DataFrame(stxt),pd.DataFrame(strrt),pd.DataFrame(stlt)],1)
```

```
# In[119]:
```

```
stack_train.shape
```

```
# In[120]:
```

```
stack = RandomForestClassifier(max_depth=13)
```

```
# In[121]:
```

```
stack.fit(stack_train,ytrain)
```

```
# In[122]:
```

```
print('\n Classifaction Report : \n\n',metrics.classification_report(ytest,stack.predict(stack_test)))
```

```
print('\n Accuracy : ',metrics.accuracy_score(ytest,stack.predict(stack_test)))
```

```
# ##### MLP
```

```
# In[123]:
```

```
from sklearn.neural_network import MLPClassifier
```

```
# In[124]:
```

```
m = MLPClassifier(learning_rate='adaptive',hidden_layer_sizes=(75,75))
```

```
# In[125]:
```

```
m.fit(xtrain,ytrain)
```

```
# In[126]:
```

```
# Classification report and accuracy
```

```
print('\n Classification Report : \n\n',metrics.classification_report(ytest,m.predict(np.matrix(xtest))))
```

```
print('Accuracy :',metrics.accuracy_score(ytest,m.predict(np.matrix(xtest))))
```

```
# ** Predictions On Test **
```

```
# In[127]:
```

```
logreg.fit(x,y)
```

```
# In[128]:
```

```
p = logreg.predict(te_x)
```

```
# In[129]:
```

```
labels = ['No','Yes']
```

```
churn=[]
```

```
for i in p:
```

```
    churn.append(labels[i])
```

```
# In[130]:
```

```
churn = pd.DataFrame(churn,columns=['Churn'])
```

```
# In[131]:
```

```
preds = pd.concat([test,churn],1) # CustomerID and Churn
```

```
# In[132]:
```

```
preds.to_csv('predictions.csv',index=False) # write to predictions.csv
```

```
# ** Patterns **
```

```
# In[133]:
```

```
exp = tree.fit(xtrain,ytrain) # gini
```

```
exp1 = entropy.fit(xtrain,ytrain) # Entropy
```

```
# In[134]:
```

```
from sklearn.tree import export_graphviz
```

```
export_graphviz(exp,'patern.dot',class_names=['No','Yes'],feature_names=xtrain.columns)
```

```
export_graphviz(exp1,'entropy.dot',class_names=['No','Yes'],feature_names=xtrain.columns)
```